

SITARA™ ARM® PROCESSORS **Boot camp**



Optimizing Linux Boot Time

This session gives an overview of methods for optimizing the boot time of a Linux system

LAB: http://processors.wiki.ti.com/index.php/Sitara_Linux_Training

July 2012

Creative Commons Attribution-ShareAlike 3.0 (CC BY-SA 3.0)



You are free:

to **Share** – to copy, distribute and transmit the work

to **Remix** – to adapt the work

to make commercial use of the work

Under the following conditions:



Attribution – You must give the original author(s) credit



Share Alike - If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.

With the understanding that:

Waiver — Any of the above conditions can be waived if you get permission from the copyright holder.

Public Domain — Where the work or any of its elements is in the public domain under applicable law, that status is in no way affected by the license.

Other Rights — In no way are any of the following rights affected by the license:

Notice — For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to this web page.



CC BY-SA 3.0 License:

<http://creativecommons.org/licenses/by-sa/3.0/us/legalcode>



Pre-work Check List

- ☐ Installed and configured VMWare Player v4 or later
- ☐ Installed Ubuntu 10.04
- ☐ Installed the latest Sitara Linux SDK and CCSv5
- ☐ Within the Sitara Linux SDK, ran the setup.sh (to install required host packages)
- ☐ Using a Sitara EVM, followed the QSG to connect ethernet, serial cables, SD card and 5V power
- ☐ Booted the EVM and noticed the Matrix GUI application launcher on the LCD
- ☐ Pulled the ipaddr of your EVM and ran remote Matrix using a web browser
- ☐ Brought the USB to Serial cable you confirmed on your setup (preferable)

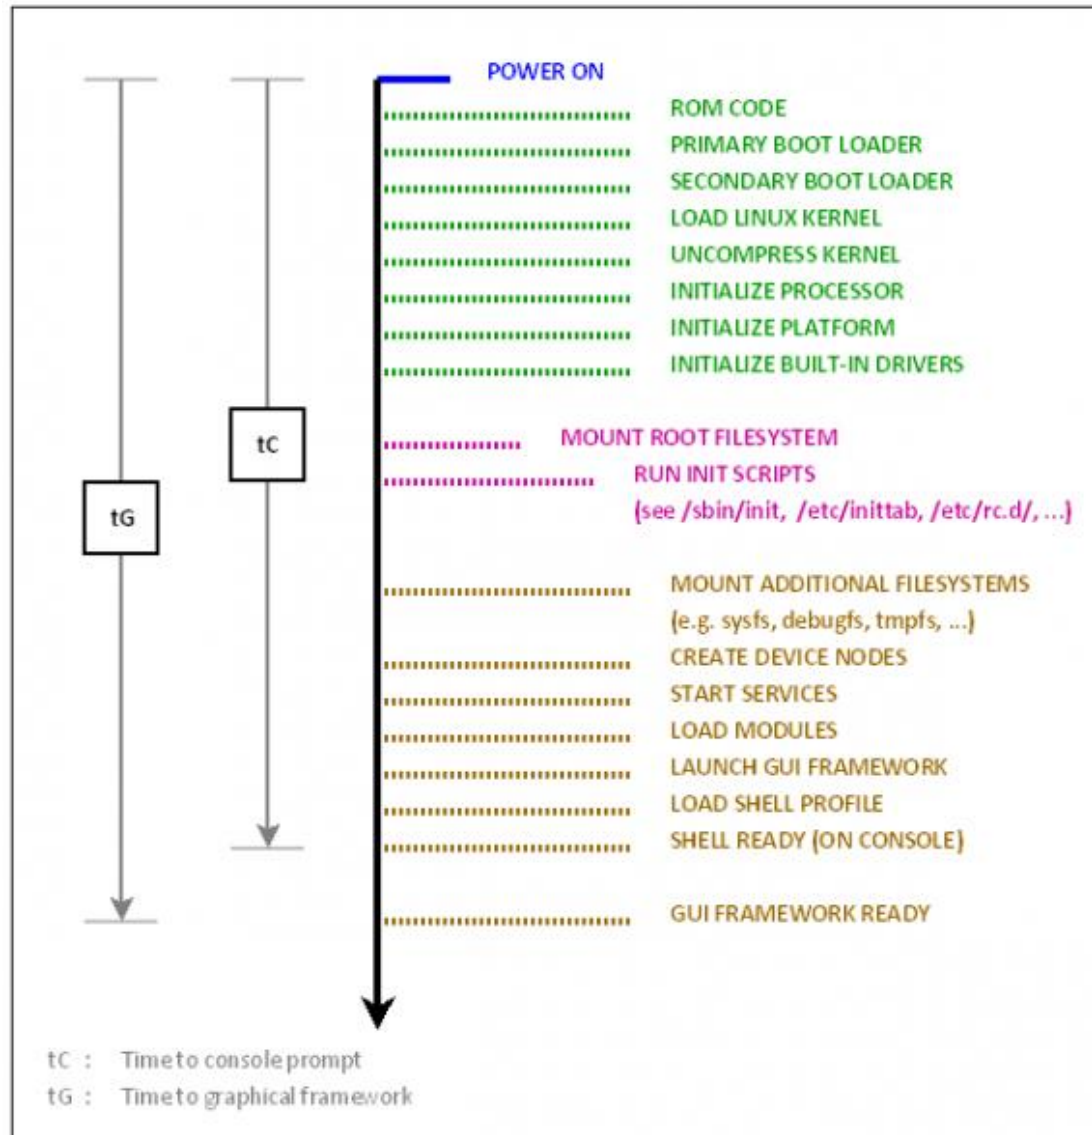
Agenda

- What is Fast Boot?
- Boot Process Overview
- Measuring Boot Time
- Identifying the Boot Steps
- Profiling
 - U-boot
 - Linux Kernel
- Optimization Techniques
 - U-boot
 - Linux Kernel
 - File System

What is Fast Boot?

- Fast boot refers to minimizing the boot time of a system. The boot time of the system is the time it takes from the application of power to the system becoming “available”
- “Available” has a lot of different meanings depending on the user expectations.
 - Appearance of the home screen for devices like cell phones
 - An audible tone or LED indicator
 - A Linux prompt
 - Becoming discoverable on the network
 - Having a key peripheral become available
- The above variances show why the Sitara Linux SDK is generally not fully optimized for boot time. Each user can have a different target and a generic SDK cannot satisfy all user targets
 - There are also many features that make for a good development environment, but which do not make for a fast booting environment

Boot Process Overview



Measuring Boot Time

- There are a couple of different options for measuring boot time. You can use a C program that time stamps each line on the serial port.
 - http://processors.wiki.ti.com/index.php/Measuring_Boot_Time
 - Compile with “gcc tstamp.c -o tstamp”
 - This can also be invoked as “cat /dev/<TTY DEVICE> | tstamp
- RealTerm for Windows supports time stamping the incoming serial data
 - <http://realterm.sourceforge.net/>
 - Support logging in Unix date format
 - Ability to log directly to a file
 - Ability to stop logging after a predefined time
 - No need to switch terminal for interactive session
- TeraTerm for Windows can now log time as well but does not give elapsed time measurements.

Measuring Boot Time - Cont

- Using the C program approach you will see output like:

```
18.066 0.093: Thu Jun 28 11:56:00 UTC 2012
INIT: Entering runlevel: 5
18.262 0.175: Starting system message bus: dbus.
18.280 0.018: Starting Hardware abstraction layer hald
19.474 1.194: Starting Dropbear SSH server: dropbear.
19.499 0.025: Starting telnet daemon.
19.538 0.039: Starting network benchmark server: netserver.
19.569 0.031: Starting syslogd/klogd: done
19.629 0.060: Starting thttpd.
19.731 0.102: Starting PVR
20.158 0.427: Starting Lighttpd Web Server: lighttpd.
20.176 0.018: 2012-06-28 11:56:02: (log.c.166) server started
```

- In the above output you can see that starting the Dropbear SSH server takes 1.194 seconds.
 - If you do not need SSH you can save 1.194 seconds by disabling Dropbear
 - Disabling Dropbear is as simple as removing the S10dropbear init script from the /etc/rc5.d directory

Identifying the Boot Steps

- One of the first things to help in optimizing boot time is to be able to recognize the markers indicating where a new part of the boot process starts
- SPL
 - The first newline character received on the serial console marks the start of SPL
- U-boot
 - The banner containing the U-boot version indicates the start of u-boot
U-Boot 2011.09 (Jun 28 2012 - 11:20:36)
- Linux Kernel
 - First line after the below line indicates the start of the Linux kernel
Uncompressing Linux... done, booting the kernel
- File System
 - The below line indicates the transition to the file system Init process
INIT: version 2.86 booting

Profiling the Boot Loaders

- There is no direct profiling support in the boot loaders, but the serial print times can still be very useful
- For example, when booting the kernel image you will see output like:

```
5.335 0.010: ## Booting kernel from Legacy Image at 80007fc0 ...  
5.339 0.004: Image Name: Arago/3.2.0-psp04.06.00.08.sdk/a  
5.344 0.005: Image Type: ARM Linux Kernel Image (uncompressed)  
5.348 0.004: Data Size: 3164688 Bytes = 3 MiB  
5.350 0.002: Load Address: 80008000  
5.352 0.002: Entry Point: 80008000  
6.273 0.921: Verifying Checksum ... OK  
6.276 0.003: XIP Kernel Image ... OK  
6.310 0.034: OK
```

- Notice that almost a second is spent verifying the kernel image checksum.
- If your system doesn't do anything about a bad image, why spend time verifying it?
 - This verification can be disabled by setting the “verify” u-boot parameter to n at the u-boot command prompt
u-boot# `setenv verify n`

Profiling the Linux Kernel

- One of the simplest ways to profile the Linux kernel is to configure “CONFIG_PRINTK_TIME” for the Linux kernel. On most recent kernels this is enabled by default.

– This adds the time since the kernel was booted in []’s before each line

```
[ 2.029785] mmc1: card claims to support voltages below the defined range. T.  
[ 2.048583] mmc1: queuing unknown CIS tuple 0x91 (3 bytes)  
[ 2.055053] mmc1: new SDIO card at address 0001  
[ 3.986724] PHY: 0:00 - Link is Up - 100/Full  
[ 4.015808] Sending DHCP requests ., OK  
[ 4.036254] IP-Config: Got DHCP answer from 0.0.0.0, my address is 128.247.10  
[ 4.044586] IP-Config: Complete:  
[ 4.047973]   device=eth0, addr=128.247.105.20, mask=255.255.254.0, gw=12,  
[ 4.056182]   host=128.247.105.20, domain=am.dhcp.ti.com, nis-domain=(non,  
[ 4.063781]   bootserver=0.0.0.0, rootserver=0.0.0.0, rootpath=
```

- In The above output we can see that it takes almost 2 seconds to get the network Phy Link up and obtain a DHCP IP address.
- Depending on network speed this could be longer. In the case where there is no DHCP server this can take minutes to timeout.
 - Setting “ip=off” in the bootargs will bypass kernel network configuration while still allowing you to configure the network in user space.

`setenv ip_method off`

Profiling the Linux Kernel - Cont

- Instrument kernel initialization. This will help you find which static drivers in the kernel are taking the most time to initialize
 - Add “initcall_debug” to the bootargs in u-boot. With recent u-boots this can usually be done using:
`setenv optargs initcall_debug`
`saveenv`
 - When the Linux system is booted you can view these initcall lines using “dmesg | grep initcall”
`[xxxxx] initcall <init function> [<module>] returned 0 after <time> usecs`
 - These lines can be sorted using the commands below to help find the modules with the largest init times
 - If CONFIG_PRINTK_TIME is enabled
`dmesg | grep initcall | sort -k8 -n`
 - If CONFIG_PRINTK_TIME is not enabled
`dmesg | grep initcall | sort -k6 -n`
 - It is possible to graphically view these initcall times using the bootgraph script in the Linux kernel sources. This requires CONFIG_PRINTK_TIME to be enabled
`cat <bootlog> | perl <kernel source dir>/scripts/bootgraph.pl > boot.svg`
- Additionally you can use other tools to help you analyze the Linux kernel such as:
 - Linux Trace Toolkit (LTTng)
 - Timing for certain kernel and process events
 - Oprofile
 - System wide profiler
 - Bootchart
 - Visualizes boot process

Areas of Optimization

- Optimizations generally fall into two areas
- Size
 - Reduce the size of binaries
 - Remove features not required to reduce component size
- Speed
 - Optimize for target processor
 - Neon optimizations
 - Use faster boot media
 - NOR/NAND vs. MMC/USB
 - Reduce the number tasks leading to boot
 - Do not check MMC if booting from NAND
 - Do not initialize network if booting from MMC
 - Reduce initialization operations
 - Do not bring up network during boot if not required
 - Do not start an SSH server on a device with no network

U-boot Optimization Techniques

- Reduce environment size so that less data is read into memory
 - CONFIG_ENV_SIZE
- Remove Unnecessary Console Print Statements
 - In board config file add

```
#define CONFIG_SILENT_CONSOLE 1
```
 - In u-boot environment do

```
setenv silent 1
```
- Set “bootdelay” to 0
- Disable un-used peripherals such as USB/MMC/Ethernet/UART
- Modify the config file for your device in <u-boot sources>/include/configs.
i.e am335x_evm.h
 - Avoid long help text for the u-boot commands to save space
 - #undef CONFIG_SYS_LONGHELP
 - Use simple parser - instead of hush
 - #undef CONFIG_SYS_HUSH_PARSER
 - If no USB/NAND/MMC/SPI/NOR(FLASH)
 - #undef CONFIG_USB *
 - #undef CONFIG_NAND
 - #undef CONFIG_MMC
 - #undef CONFIG_SPI
 - #undef CONFIG_FLASH_*

U-boot Optimization Techniques - Cont

- Remove -g option from the compiler
- Other Ideas
 - Disable UART boot
 - Remove Image Verification (covered before)
 - Perhaps try uncompressed image
 - Verify that kernel image is read to the proper memory location

Linux Optimization Techniques

- Remove un-necessary drivers/features from kernel configuration
 - Reduces driver initialization time
 - Reduces kernel size
- Build non-fast boot drivers as modules
 - Load them after the system is booted when there is more time
- Disable console output using “quiet”
 - Displaying messages on console takes time
 - Setting the “quiet” option in the bootargs disables display on console but messages are still logged
 - `u-boot# setenv optargs quiet`
 - Remove un-used consoles. These take time to initialize.
 - These can be removed in the /etc/inittab file on the target file system
 - It is possible to completely disable printk but this will eliminate a lot of debug information

Linux Optimization Techniques - Cont

- Defer module init calls
 - It is possible to defer module init calls without having to build the modules dynamically.
 - This requires modifying the kernel
 - For modules that are not needed at boot replace the “module_init()” function calls to “deferred_module_init()”
 - Once the system booted the deferred calls can be executed by doing:
`echo 1 > /proc/deferred_initcalls`
 - Additional details at http://elinux.org/Deferred_Initcalls
- Remove the -g option from the compile
- Disable kernel debugging features
 - Kernel debugging
 - Debug Filesystem (NOTE: Some features may need this)
 - Tracers
- Remove any instrumentation you may have added such as initcall instrumentation.

Linux Optimization Techniques - Cont

- Pre-set loops per jiffy
 - You just need to measure this once
 - Find lpj value in kernel boot messages
 - Calibrating delay loop... 718.02 BogoMIPS (lpj=3590144)
 - Add the “lpj=3590144” to the bootargs
- Use Static IP addressing where possible
 - If you don't need networking then disable it altogether
 - If you want networking capability but not NFS then be sure to set “ip=off” on the kernel command line.
- Set memory limit with “mem=” option
 - Use only as much memory as needed to avoid DDR initialization time

File System Optimization Techniques

- Use minimal BusyBox file system
 - Reduces forking in shell
 - Build static if possible to reduce the need for un-used code in the file system
 - Be careful because static linking can also cause your file system size to increase dramatically.
- Avoid using ramdisk or initramfs
 - Must load entire ramdisk from flash into DDR
 - May only need a small part at boot time
 - i.e. may not need all of glibc but entire library will be loaded into DDR using ramdisk or initramfs
 - Buffer cache can keep frequently used files in memory

File System Optimization Techniques - Cont

- Pre-linking
 - Avoid run time linking penalty
 - Drawback: If library changes app must be rebuilt.
- Use tmpfs file system
 - No need to initialize file system for non-persistent data
- Use split file systems
 - Have multiple file system partitions
 - Put only the files needed for boot in the root file system partition
 - Put other files in second file system which can be mounted after boot

File System Optimization Techniques - Cont

- Strip executables
 - Removes un-needed symbols and reduces size
- Avoid udev for static systems
 - If the system doesn't change then make device nodes manually rather than udev creating them
 - Hotplug-daemon can still run later to add additional devices that are plugged in
- Disable init scripts that start unneeded services
- Use GNU_HASH to make dynamic linking faster (This is the default for the SDK)
- Use systemd for parallel initialization

Credits/Sources

- http://processors.wiki.ti.com/index.php/Optimize_Linux_Boot_Time
 - by Sanjeev Premi
- http://elinux.org/Boot_Time

SITARA™ ARM® PROCESSORS

Boot camp



For more Sitara Boot Camp sessions visit:
www.ti.com/sitarabootcamp

THANK YOU!